


Non-Pipelined Processors

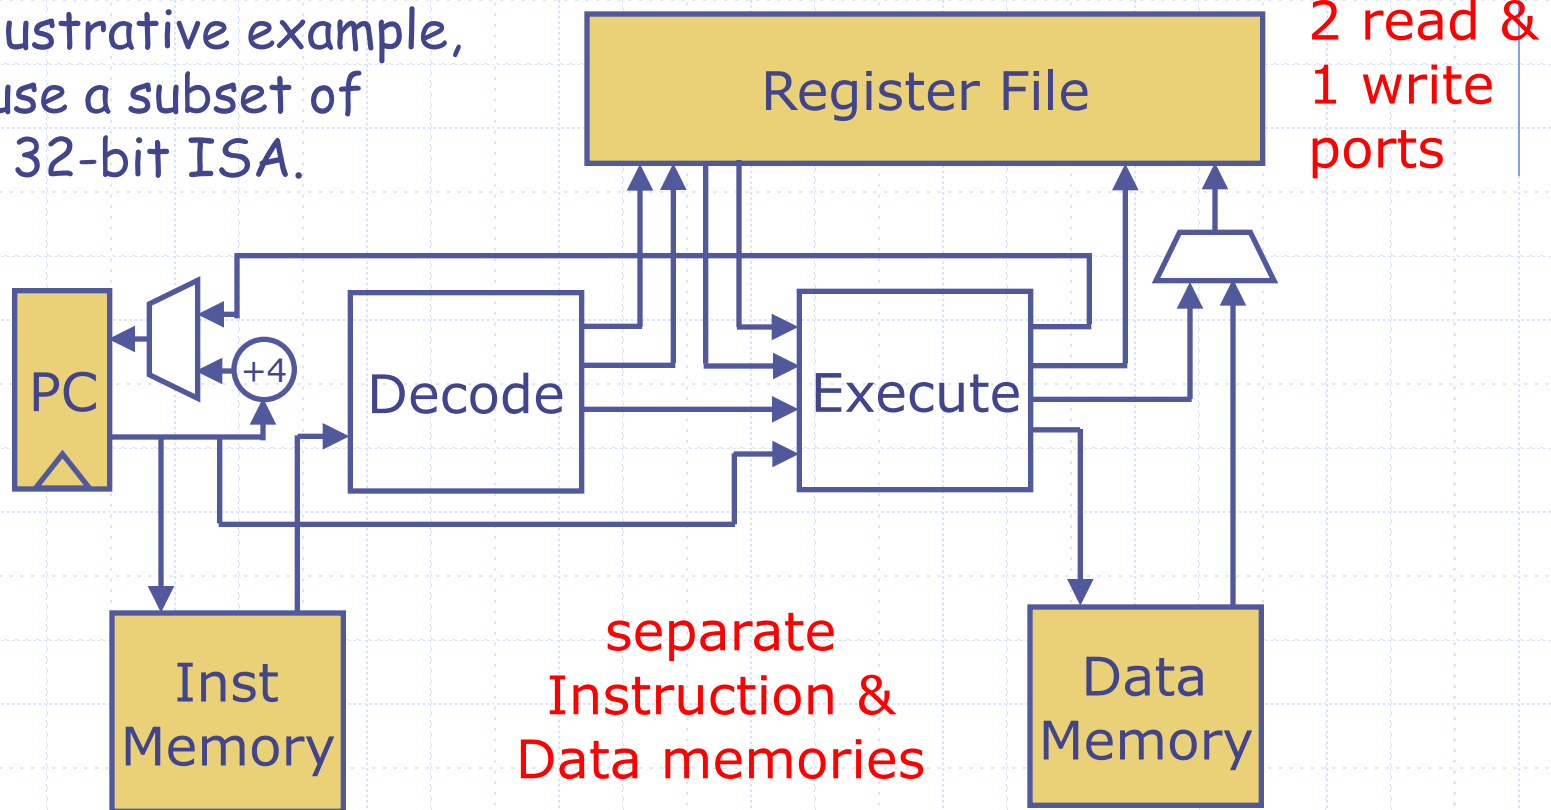
Arvind
Computer Science & Artificial Intelligence Lab.
Massachusetts Institute of Technology

The Big Picture



Single-Cycle RISC Processor

As an illustrative example, we will use a subset of RISC-V 32-bit ISA.



Datapath and control are derived automatically from a high-level rule-based description

Single-Cycle Implementation

code structure

```
module mkProc (Proc);
```

```
  Reg# (Addr)  pc <- mkRegU;
```

```
  RFile        rf <- mkRFile;
```

```
  IMemory      iMem <- mkIMemory;
```

```
  DMemory      dMem <- mkDMemory;
```

to be explained later

instantiate the state

```
rule doProc;
```

```
  let inst = iMem.req(pc);
```

```
  let dInst = decode(inst);
```

```
  let rVal1 = rf.rd1(dInst.rSrc1);
```

```
  let rVal2 = rf.rd2(dInst.rSrc2);
```

```
  let eInst = exec(dInst, rVal1, rVal2, pc);
```

```
  update rf, pc and dMem
```

extracts fields
needed for
execution

produces values
needed to
update the
processor state

RISC-V Register States

◆ 32 general purpose registers (GPR)

- x0, x1, ..., x31
- 32-bit wide integer registers
- x0 is hard-wired to zero

◆ Program counter (PC)

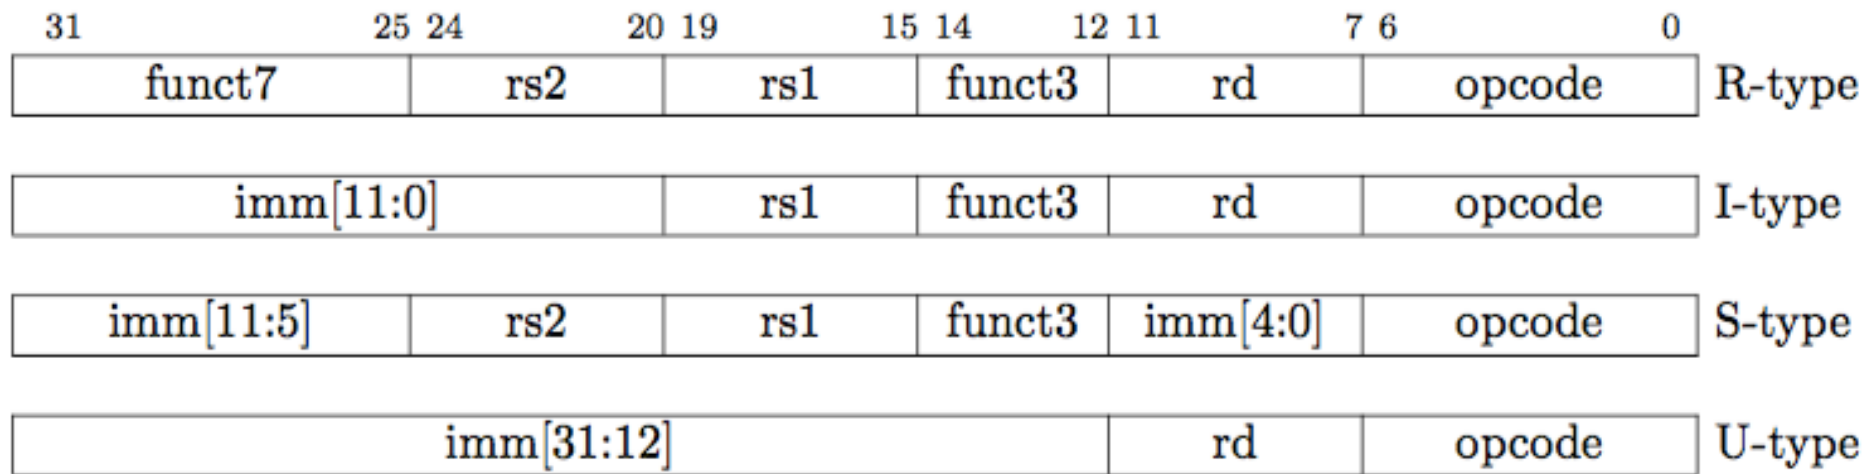
- 32-bit wide

◆ CSR (Control and Status Registers)

- cycle
- instret
- mhartid
- mtohost
- ...

} will be
implemented
in labs as
needed

Instruction Formats



Computational Instructions

◆ Register-Register instructions (R-type)



- opcode=OP: $rd \leftarrow rs1 \text{ (funct3, funct7) } rs2$
- funct3 = SLT/SLTU/AND/OR/XOR/SLL
- funct3= ADD
 - ◆ funct7 = 0000000: $rs1 + rs2$
 - ◆ funct7 = 0100000: $rs1 - rs2$
- funct3 = SRL
 - ◆ funct7 = 0000000: logical shift right
 - ◆ funct7 = 0100000: arithmetic shift right

Computational Instructions *cont*

◆ Register-immediate instructions (I-type)



- opcode = OP-IMM: $rd \leftarrow rs1$ (funct3) I-imm
- I-imm = signExtend(inst[31:20])
- funct3 = ADDI/SLTI/SLTIU/ANDI/ORI/XORI

◆ A slight variant in coding for shift instructions - SLLI / SRLI / SRAI


- $rd \leftarrow rs1$ (funct3, inst[30]) I-imm[4:0]

Shift Instructions

31	25 24	20 19	15 14	12 11	7 6	0
imm[11:5]	imm[4:0]	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
0000000	shamt[4:0]	src	SLLI	dest	OP-IMM	
0000000	shamt[4:0]	src	SRLI	dest	OP-IMM	
0100000	shamt[4:0]	src	SRAI	dest	OP-IMM	

Control Instructions

◆ Unconditional jump and link (UJ-type)





- opcode = JAL: $rd \leftarrow pc + 4$; $pc \leftarrow pc + J-imm$
- J-imm = $signExtend(\{inst[31], inst[19:12], inst[20], inst[30:21], 1'b0\})$
- Jump $\pm 1MB$ range

◆ Unconditional jump via register and link (I-type)



- opcode = JALR: $rd \leftarrow pc + 4$; $pc \leftarrow (rs1 + I-imm) \& \sim 0x01$
- I-imm = $signExtend(inst[31:20])$

Control Instructions



Control Instructions *cont.*

◆ Conditional branches (SB-type)



- opcode = BRANCH: $pc \leftarrow \text{compare}(\text{funct3}, rs1, rs2) ? pc + B\text{-imm} : pc + 4$
- B-imm = $\text{signExtend}(\{\text{inst}[31], \text{inst}[7], \text{inst}[30:25], \text{inst}[11:8], 1'b0\})$
- Jump $\pm 4\text{KB}$ range
- funct3 = BEQ/BNE/BLT/BLTU/BGE/BGEU

1'b0 means it's half-word aligned. This is because RISC V allows 16-bit compressed format of instructions

Control Instructions

31	30	25 24	20 19	15 14	12 11	8	7	6	0
imm[12]	imm[10:5]	rs2	rs1	funct3	imm[4:1]	imm[11]	opcode		
1	6	5	5	3	4	1	7		
offset[12,10:5]		src2	src1	BEQ/BNE	offset[11,4:1]		BRANCH		
offset[12,10:5]		src2	src1	BLT[U]	offset[11,4:1]		BRANCH		
offset[12,10:5]		src2	src1	BGE[U]	offset[11,4:1]		BRANCH		

Load & Store Instructions

◆ Load (I-type)





- opcode = LOAD: $rd \leftarrow \text{mem}[rs1 + \text{I-imm}]$
- I-imm = $\text{signExtend}(\text{inst}[31:20])$
- funct3 = LW/LB/LBU/LH/LHU

◆ Store (S-type)



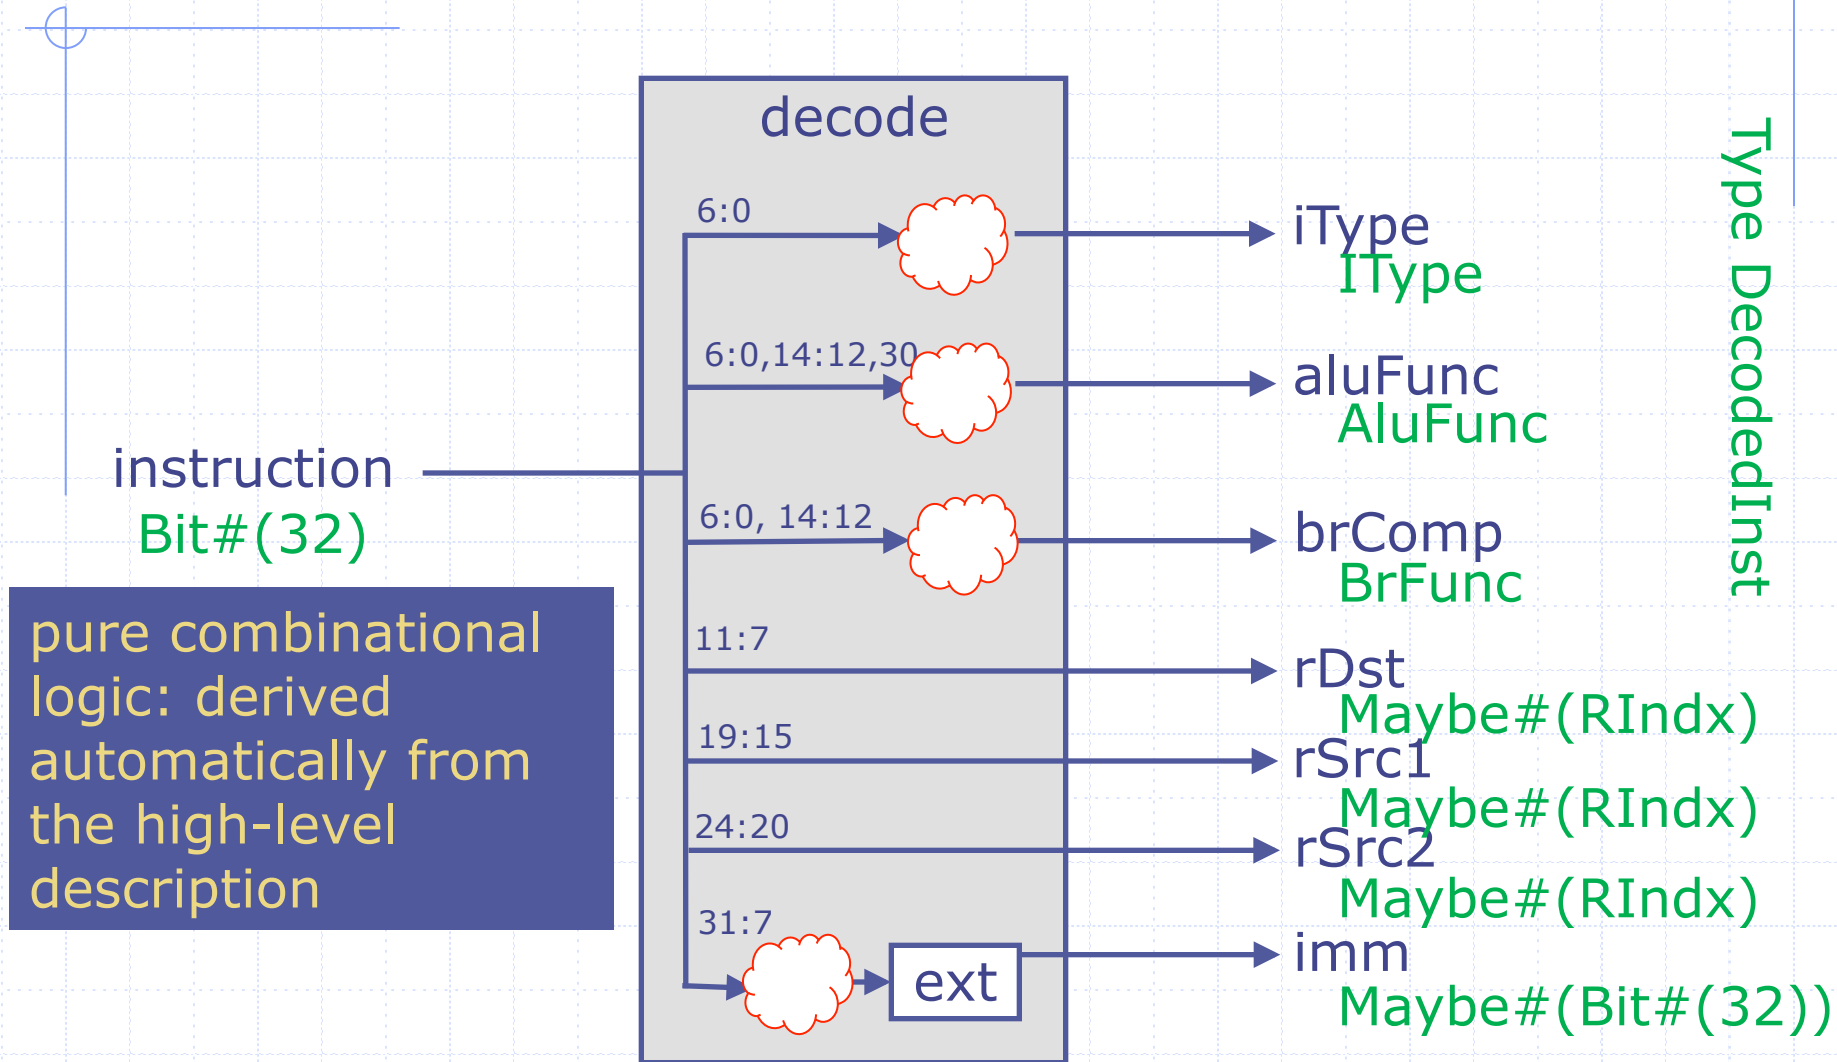
- opcode = STORE: $\text{mem}[rs1 + \text{S-imm}] \leftarrow rs2$
- S-imm = $\text{signExtend}(\{\text{inst}[31:25], \text{inst}[11:7]\})$
- funct3 = SW/SB/SH

Load and Store Instructions



Decoding Instructions:

extract fields needed for execution



Decoded Instruction Type

```
typedef struct {  
    IType          iType;  
    AluFunc         aluFunc;  
    BrFunc          brFunc;  
    Maybe#(RIdx)    dst;  
    Maybe#(RIdx)    src1;  
    Maybe#(RIdx)    src2;  
    Maybe#(Data)    imm;  
} DecodedInst deriving (Bits, Eq);  
  
typedef enum {Unsupported, Alu, Ld, St, J, Jr, Br, Auipc}  
IType deriving (Bits, Eq);  
  
typedef enum {Add, Sub, And, Or, Xor, Slt, Sltu, Sll,  
Sra, Srl} AluFunc deriving (Bits, Eq);  
  
typedef enum {Eq, Neq, Lt, Ltu, Ge, Geu, AT, NT} BrFunc  
deriving (Bits, Eq);
```

Destination register 0 behaves like an Invalid destination

Instruction groups with similar executions paths

Internal names for various opcode and funct3 patterns

// opcode

Bit#(7) opOpImm = 7'b0010011; // OP-IMM

Bit#(7) opOp = 7'b0110011; // OP

Bit#(7) opLui = 7'b0110111; // LUI

Bit#(7) opAuipc = 7'b0010111; // AUIPC

Bit#(7) opJal = 7'b1101111; // JAL

Bit#(7) opJalr = 7'b1100111; // JALR

Bit#(7) opBranch = 7'b1100011; // BRANCH

Bit#(7) opLoad = 7'b0000011; // LOAD

Bit#(7) opStore = 7'b0100011; // STORE

// funct3

Bit#(3) fnADD = 3'b000; // ADD

Bit#(3) fnSLL = 3'b001; // SLL

Bit#(3) fnSLT = 3'b010; // SLT

Values are
specified in
the RISC-V
ISA

Decode Function

```
function DecodedInst decode (Bit#(32) inst);  
  DecodedInst dInst = ?;  
  let opcode = inst[ 6 : 0 ];  
  let rd     = inst[ 11 : 7 ];  
  let funct3 = inst[ 14 : 12 ];  
  let rs1    = inst[ 19 : 15 ];  
  let rs2    = inst[ 24 : 20 ];  
  let aluSel = inst[ 30 ]; // Add/Sub, Srl/Sra  
  Bit#(32) immI=...; Bit#(32) immS=...; Bit#(32) immB=...;  
  Bit#(32) immU=...; Bit#(32) immJ=...; // I/S/B/U/J-imm  
  case (opcode)  
    opOp  
    ...  
  endcase  
  return dInst;  
endfunction
```

initially undefined

Decoding Instructions: Computational Instructions

```
opOp: begin
    dInst.iType = Alu;
    dInst.aluFunc = case (funct3)
        fnAND: And;
        fnSLTU: Sltu;
        ...
        fnADD: Add;
        fnSR: aluSel == 0 ? Srl : Sra;
    endcase;
    dInst.brFunc = NT;
    dInst.dst = Valid rd;
    dInst.src1 = Valid rs1;
    dInst.src2 = Invalid;
    dInst.imm = Valid ImmI;
end
```

Decoding instructions
with immediate operand
(i.e. opcode = OP-IMM) is
similar

Decoding Instructions: Conditional Branch

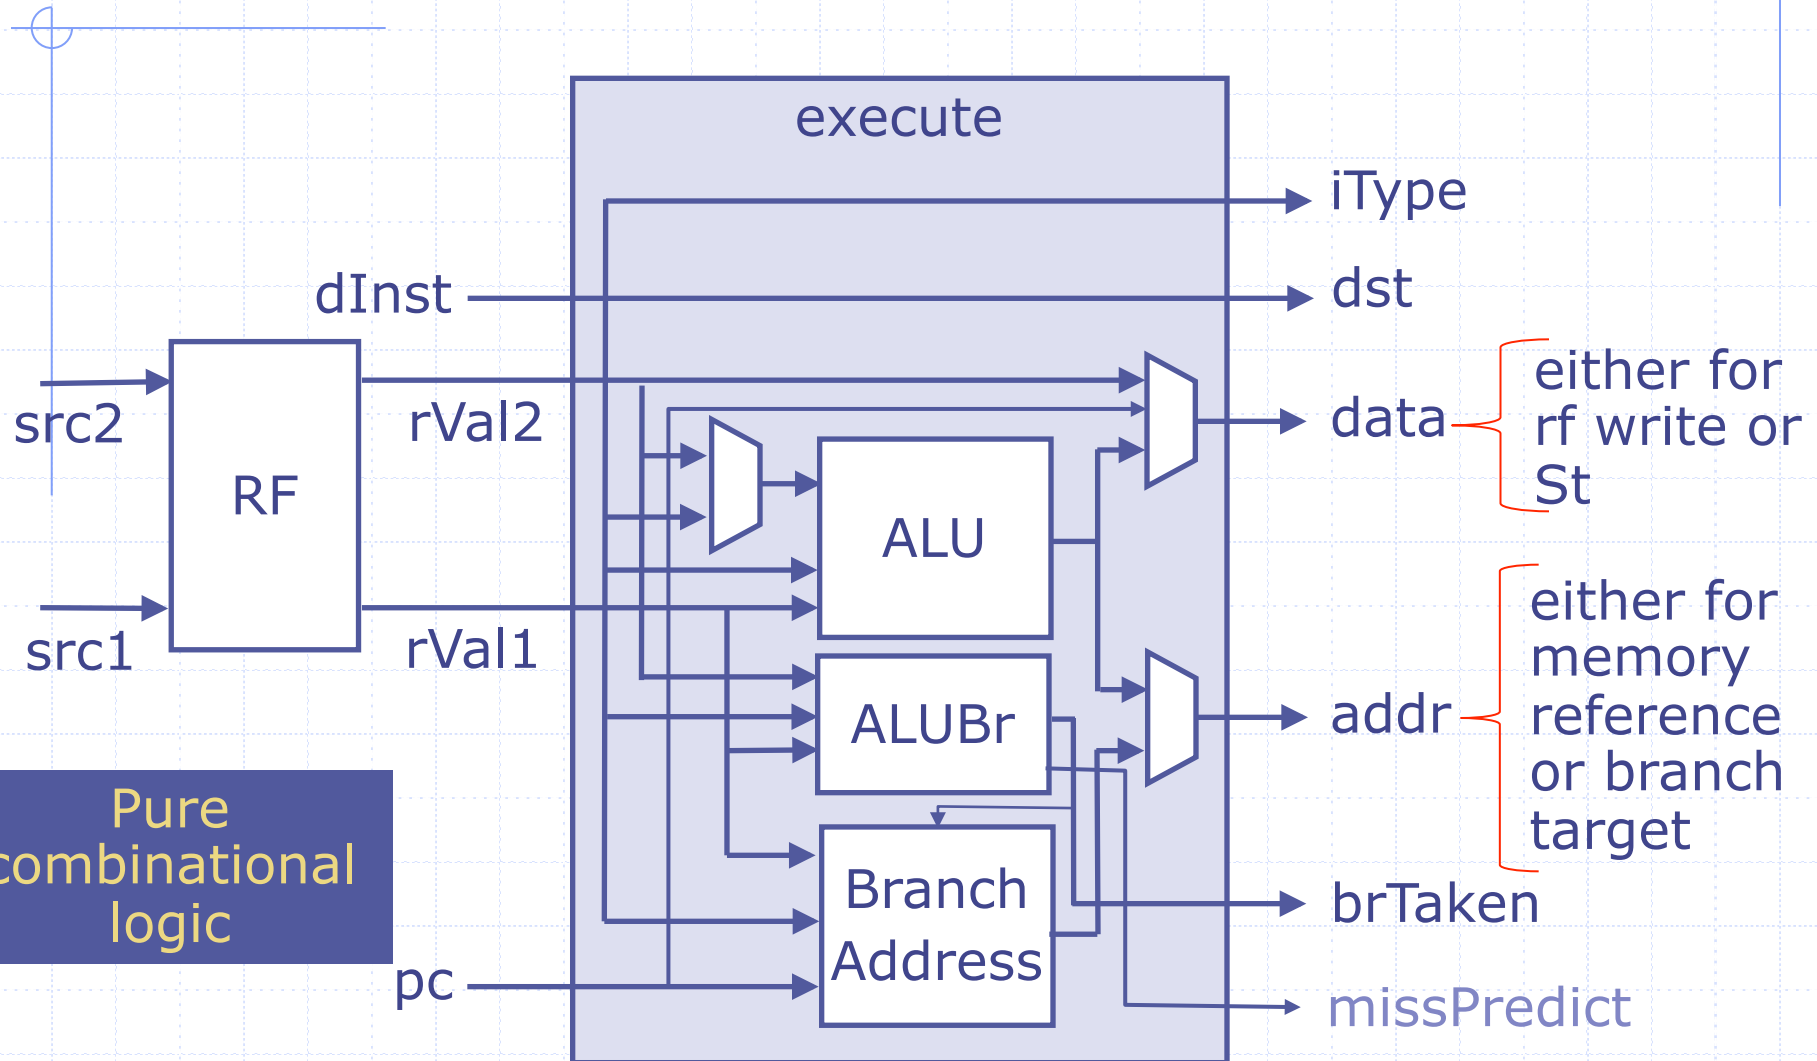
```
opBranch: begin
  Maybe#(BrFunc) brF =
    case (funct3)
    fnBEQ: Valid Eq;
    ...
    fnBGEU: Valid Geu;
    default: Invalid;
  endcase;
  dInst.iType = isValid(brF) ? Br : Unsupported;
  dInst.aluFunc = ?;
  dInst.brFunc = fromMaybe(?, brF);
  dInst.dst = Invalid;
  dInst.src1 = Valid rs1;
  dInst.src2 = Valid rs2;
  dInst.imm = Valid immB;
end
```

Decoding Instructions: Load & Store

```
opLoad: begin // only support LW
    dInst.iType = funct3 == fnLW ? Ld : Unsupported;
    dInst.aluFunc = Add; // calc effective addr
    dInst.brFunc = NT;
    dInst.dst = Valid rd;
    dInst.src1 = Valid rs1;
    dInst.src2 = Invalid;
    dInst.imm = Valid immI;
end

opStore: begin // only support SW
    dInst.iType = funct3 == fnSW ? St : Unsupported;
    dInst.aluFunc = Add; // calc effective addr
    dInst.brFunc = NT;
    dInst.dst = Invalid;
    dInst.src1 = Valid rs1;
    dInst.src2 = Valid rs2;
    dInst.imm = Valid immS;
end
```

Reading Registers and Executing Instructions



Output type of exec function

```
typedef struct {  
    IType          iType;  
    Maybe# (RIndx) dst;  
    Data           data;  
    Addr           addr;  
    Bool           mispredict;  
    Bool           brTaken;  
} ExecInst deriving (Bits, Eq);
```


Execute Function

```
function ExecInst exec(DecodedInst dInst, Data rVal1,
                        Data rVal2, Addr pc);
    ExecInst eInst = ?;
    Data aluVal2 = fromMaybe(rVal2, dInst.imm);

    let aluRes = alu(rVal1, aluVal2, dInst.aluFunc);
    eInst.iType = dInst.iType;
    eInst.data = case (dInst.iType)
                St : rVal2;
                J, Jr: (pc+4);
                Auipc: (pc+fromMaybe(?, dInst.imm));
                default: aluRes; endcase

    let brTaken = aluBr(rVal1, rVal2, dInst.brFunc);
    let brAddr = brAddrCalc(pc, rVal1, dInst.iType,
                           fromMaybe(?, dInst.imm), brTaken);

    eInst.brTaken = brTaken;
    eInst.addr = (dInst.iType==Ld || dInst.iType==St)?
                 aluRes : brAddr;

    eInst.dst = dInst.dst;
    return eInst;
endfunction
```

Needed to load PC into a register

Single-Cycle SMIPS *atomic state updates*

```
if (eInst.iType == Ld)
  eInst.data <- dMem.req(MemReq{op: Ld,
                               addr: eInst.addr, data: ?});
else if (eInst.iType == St)
  let dummy <- dMem.req(MemReq{op: St,
                               addr: eInst.addr, data: data});

if (isValid(eInst.dst))
  rf.wr(fromMaybe(?, eInst.dst), eInst.data);

pc <= eInst.brTaken ? eInst.addr : pc + 4;
```

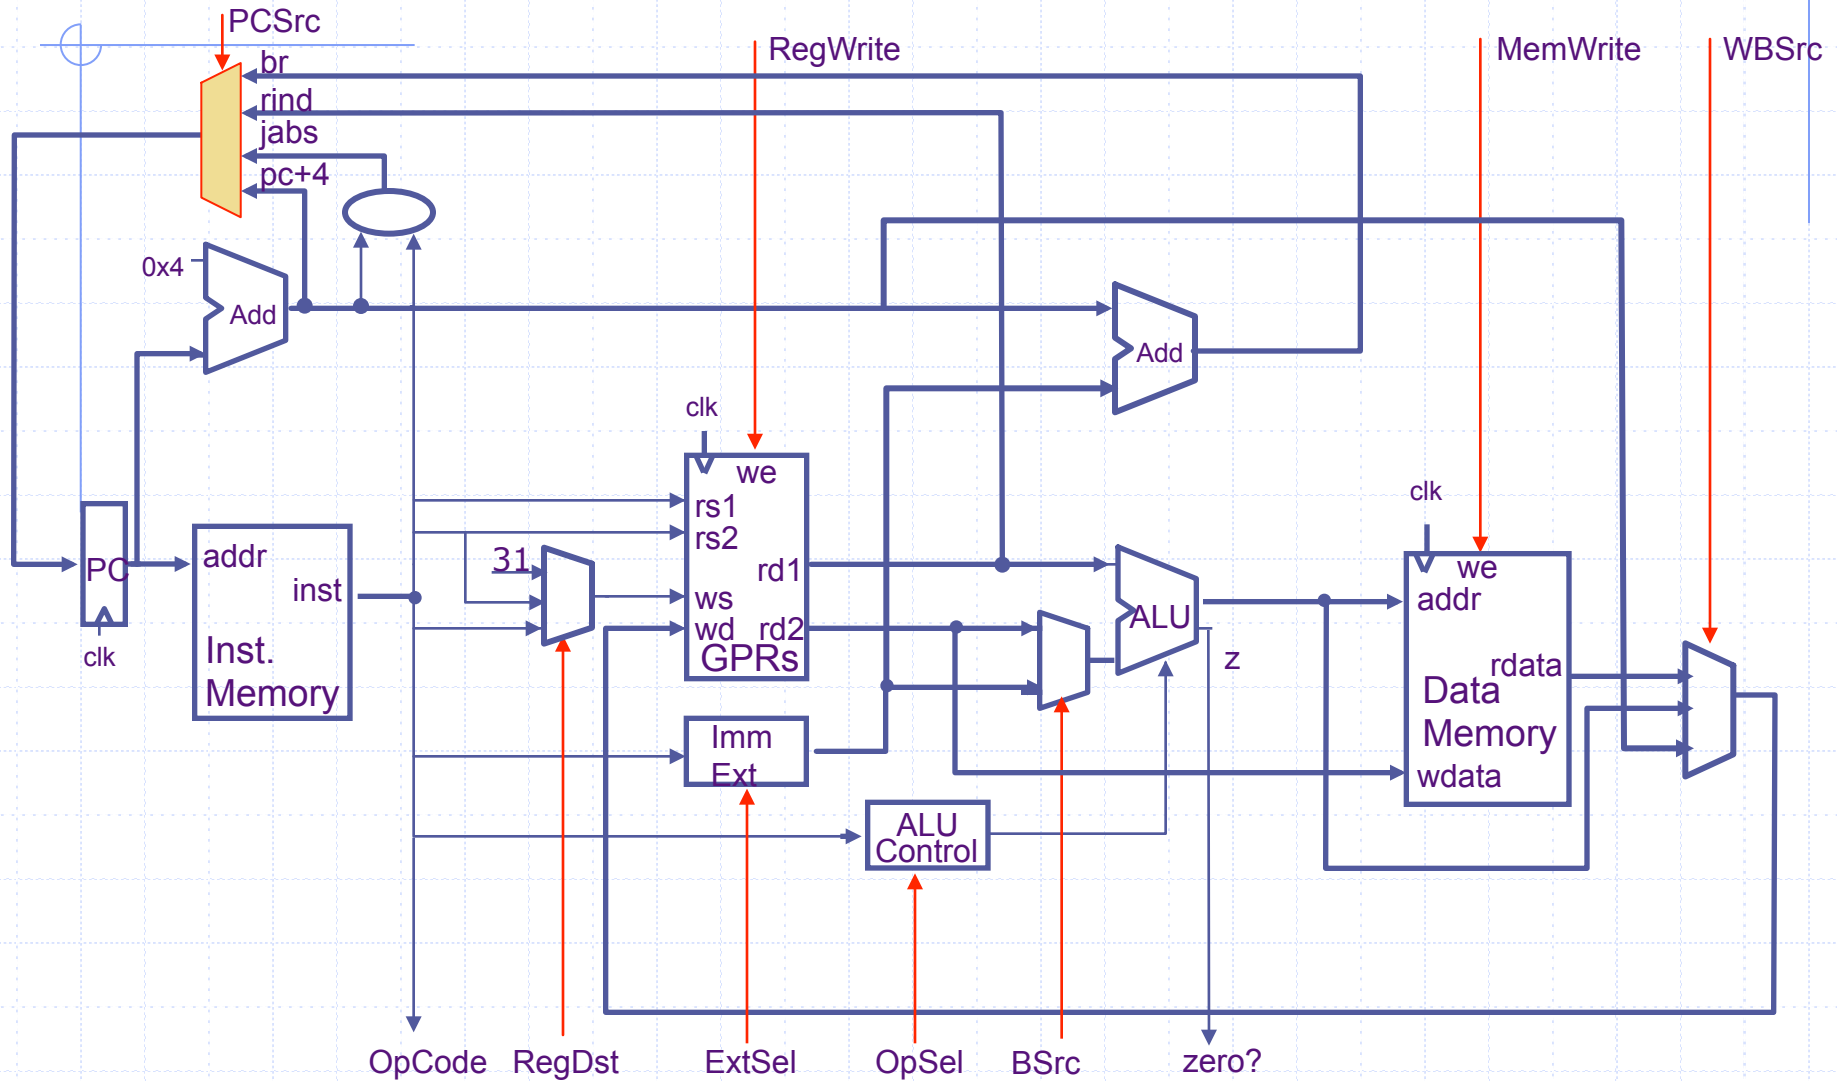
endrule
endmodule

state updates

The whole processor is described using one rule;
lots of big combinational functions

Harvard-Style Datapath for MIPS

old way



Hardwired Control Table

Opcode	ExtSel	BSrc	OpSel	MemW	RegW	WBSrc	RegDst	PCSrc
ALU	*	Reg	Func	no	yes	ALU	rd	pc+4
ALUi	sExt ₁₆	Imm	Op	no	yes	ALU	rt	pc+4
ALUiu	uExt ₁₆	Imm	Op	no	yes	ALU	rt	pc+4
LW	sExt ₁₆	Imm	+	no	yes	Mem	rt	pc+4
SW	sExt ₁₆	Imm	+	yes	no	*	*	pc+4
BEQZ _{z=0}	sExt ₁₆	*	0?	no	no	*	*	br
BEQZ _{z=1}	sExt ₁₆	*	0?	no	no	*	*	pc+4
J	*	*	*	no	no	*	*	jabs
JAL	*	*	*	no	yes	PC	R31	jabs
JR	*	*	*	no	no	*	*	rind
JALR	*	*	*	no	yes	PC	R31	rind

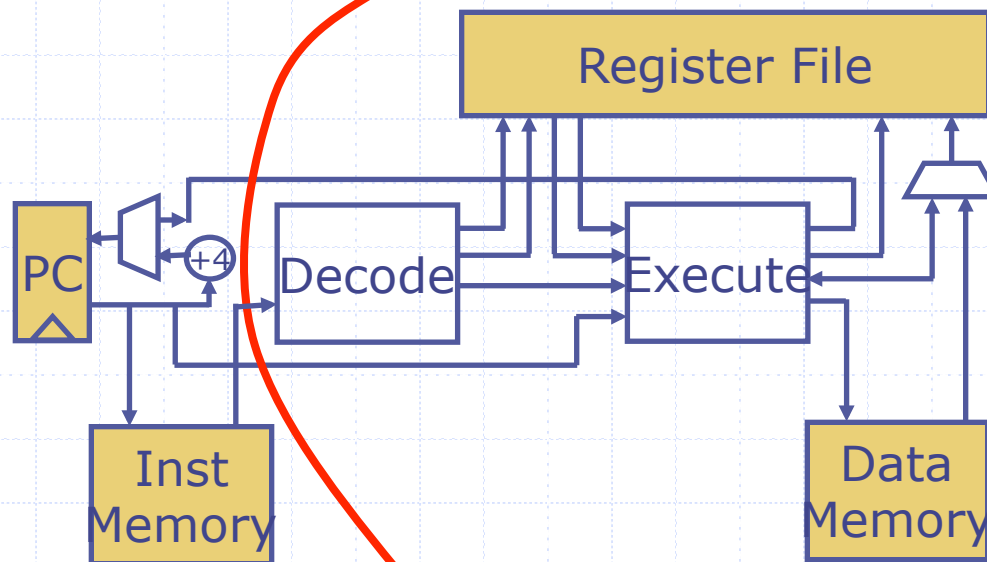
BSrc = Reg / Imm

RegDst = rt / rd / R31

WBSrc = ALU / Mem / PC

PCSrc = pc+4 / br / rind / jabs

Single-Cycle RISC-V: *Clock Speed*



$$t_{\text{Clock}} > t_M + t_{\text{DEC}} + t_{\text{RF}} + t_{\text{ALU}} + t_M + t_{\text{WB}}$$

We can improve the clock speed if we execute each instruction in two clock cycles

$$t_{\text{Clock}} > \max \{t_M, (t_{\text{DEC}} + t_{\text{RF}} + t_{\text{ALU}} + t_M + t_{\text{WB}})\}$$

However, this may not improve the performance because each instruction will now take two cycles to execute

Structural Hazards

- ◆ Sometimes multicycle implementations are necessary because of resource conflicts, aka, *structural hazards*
 - Princeton style architectures use the same memory for instruction and data and consequently, require at least two cycles to execute Load/Store instructions
 - If the register file supported less than 2 reads and one write concurrently then most instructions would take more than one cycle to execute
- ◆ Usually extra registers are required to hold values between cycles

Extras Slides on Decoding

Instruction Formats

◆ R-type instruction




◆ I-type instruction & I-immediate (32 bits)



I-imm = signExtend(inst[31:20])

◆ S-type instruction & S-immediate (32 bits)



S-imm = signExtend({inst[31:25], inst[11:7]})

Instruction Formats *cont.*

◆ SB-type instruction & B-immediate (32 bits)




B-imm = signExtend({inst[31], inst[7], inst[30:25], inst[11:8], 1'b0})

◆ U-type instruction & U-immediate (32 bits)



U-imm = signExtend({inst[31:12], 12'b0})

◆ UJ-type instruction & J-immediate (32 bits)



J-imm = signExtend({inst[31], inst[19:12], inst[20], inst[30:21], 1'b0})

Computational Instructions *cont.*

◆ Register-immediate instructions (U-type)



- opcode = LUI : $rd \leftarrow U-imm$
- opcode = AUIPC : $rd \leftarrow pc + U-imm$
- $U-imm = \{inst[31:12], 12'b0\}$

Decoding Instructions

```
case (opcode)
  opOpImm: ...
  opOp: ...
  opLui: ...
  opAuipc: ...
  opJal: ...
  opJalr: ...
  opBranch: ...
  opLoad: ...
  opStore: ...
  default: ... // Unsupported
endcase;
```

```
// opcode
Bit#(7) opOpImm = 7'b0010011;
Bit#(7) opOp    = 7'b0110011;
Bit#(7) opLui   = 7'b0110111;
Bit#(7) opAuipc = 7'b0010111;
Bit#(7) opJal   = 7'b1101111;
Bit#(7) opJalr  = 7'b1100111;
Bit#(7) opBranch = 7'b1100011;
Bit#(7) opLoad  = 7'b0000011;
Bit#(7) opStore = 7'b0100011;
```

Decoding Instructions:

Computational Instructions - Imm

```
opOpImm: begin
    dInst.iType = Alu;
    dInst.aluFunc = case (funct3)
        fnADD:    Add;
        fnSLTU:   Sltu;
        ...
        fnSLL:    Sll;
        fnSR:     aluSel == 0 ? Srl : Sra;
    endcase;
    dInst.brFunc = NT;
    dInst.dst = Valid rd;
    dInst.src1 = Valid rs1;
    dInst.src2 = Invalid;
    dInst.imm = Valid immI;
end
```

Decoding Instructions: Computational Instructions *cont.*

```
opLui: begin // rd = immU + r0
    dInst.iType = Alu;
    dInst.aluFunc = Add;
    dInst.brFunc = NT;
    dInst.dst = tagged Valid rd;
    dInst.src1 = tagged Valid 0;
    dInst.src2 = tagged Invalid;
    dInst.imm = tagged Valid immU;
```

end

```
opAuipc: begin
    dInst.iType = Auipc;
    dInst.aluFunc = ?;
    dInst.brFunc = NT;
    dInst.dst = tagged Valid rd;
    dInst.src1 = tagged Invalid;
    dInst.src2 = tagged Invalid;
    dInst.imm = tagged Valid immU;
```

end

Decoding Instructions: Unconditional Jumps

```
opJal: begin
    dInst.iType = J;
    dInst.aluFunc = ?;
    dInst.brFunc = AT;
    dInst.dst = Valid rd;
    dInst.src1 = Invalid;
    dInst.src2 = Invalid;
    dInst.imm = Valid immJ;
end
opJalr: begin
    dInst.iType = Jr;
    dInst.aluFunc = ?;
    dInst.brFunc = AT;
    dInst.dst = Valid rd;
    dInst.src1 = Valid rs1;
    dInst.src2 = Invalid;
    dInst.imm = Valid immI;
end
```

Decoding instructions: Unsupported

```
default: begin  
    dInst.iType = Unsupported;  
    dInst.aluFunc = ?;  
    dInst.brFunc = NT;  
    dInst.dst = Invalid;  
    dInst.src1 = Invalid;  
    dInst.src2 = Invalid;  
    dInst.imm = Invalid;  
end
```

Branch Address Calculation

```
function Addr brAddrCalc(Addr pc, Data val,  
                          IType iType, Data imm, Bool taken);  
  Addr pcPlus4 = pc + 4;  
  Addr targetAddr = case (iType)  
    J   : {pc + imm};  
    Jr  : {truncateLSB(val + imm), 1'b0};  
    Br  : (taken ? pc + imm : pcPlus4);  
    default: pcPlus4;  
  endcase;  
  return targetAddr;  
endfunction
```